

A Temporal Fusion Transformer (TFT) Approach to Sector Rotation



*“Oh, those trophies? Father says we
come from old GameStop money.”*

Abstract

This paper presents a robust framework for weekly sector rotation using a Temporal Fusion Transformer (TFT)¹. Unlike standard forecasting models that predict absolute price, this architecture is optimized to predict relative returns (Sector - SPY) to isolate idiosyncratic sector strength from broad market movements. We implement a

¹ Lim, Bryan, et al. "Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting." *International Journal of Forecasting*, vol. 37, no. 4, 2021, pp. 1748-1764.
<https://doi.org/10.1016/j.ijforecast.2021.03.012>.

walk-forward validation schema and a multi-seed ensembling approach to mitigate the high variance inherent in financial time-series data.

I. Model Architecture

The core engine is a modified TFT designed for interpretability and feature selection. The model architecture is specifically designed to handle the non-stationary nature of financial markets by incorporating several key components:

- **Variable Selection Network (VSN):** Learns which features—such as 4-week RSI, realized volatility, or relative volume—carry the most predictive weight for specific sectors at any given timestep. This allows the model to dynamically ignore noisy signals.
- **Encoder-Decoder Temporal Logic:** Employs an LSTM encoder to capture sequential momentum and a self-attention mechanism to identify critical historical weeks within a 12-week lookback period.
- **Output Head:** Instead of a single regression point, the model employs per-sector output projections to generate a full cross-sectional ranking of the 11 S&P 500 sectors.

II. Feature Engineering

The model ingests thirteen features per variate at each timestep, producing an input tensor of shape (lookback length $\times$ 18 variates $\times$ 13 features) for each weekly observation. The eighteen variates comprise the eleven sector ETFs, the SPY benchmark, and six macroeconomic variates (Treasury, credit, dollar, and volatility instruments). The feature set was designed to provide diverse signal types—directional, volatility-based, volume-based, short- and multi-horizon momentum, and cross-sectional—so that the Variable Selection Network has meaningful candidates to weight differently across sectors and market regimes.

- **Returns (returns, log_returns, open_close_return):** Weekly close-to-close simple and log returns are both included. Simple returns are the natural unit for portfolio mathematics; log returns are statistically better behaved and more nearly normally distributed, which benefits gradient-based training. Open-to-close return is also computed to capture intra-week directional pressure independent of weekend gap effects.
- **Volatility (weekly_range, realized_vol_4):** Weekly high-low range and 4-week realized volatility provide complementary risk signals. Range captures immediate uncertainty; realized volatility identifies the prevailing volatility regime. Sector rotation patterns differ markedly between high-volatility and low-volatility environments, and explicit volatility features help the model distinguish them.

- **Volume (log_volume, relative_volume):** Log-transformed volume handles the wide dynamic range of trading activity. Relative volume (current volume divided by a recent baseline) flags unusual activity. Volume confirmation is a long-established quantitative finance principle: price moves accompanied by elevated volume tend to be more durable than moves on thin volume, and relative-volume spikes often precede sustained trends.
- **Momentum (rsi_4, price_vs_sma4):** 4-week Relative Strength Index and price-versus-4-week-SMA (Simple Moving Average) both quantify recent directional strength. RSI provides a bounded oscillator sensitive to overbought and oversold conditions; the SMA ratio provides a continuous trend indicator. The 4-week window is short enough to react to regime shifts but long enough to filter out single-week noise.
- **Cross-Sectional Rank (return_rank):** The rank of each sector's recent return relative to the other ten sectors is included as a direct input. This feature is the most closely matched to the prediction target: because the model is predicting relative outperformance, providing the current relative-performance ordering accelerates learning and reduces the burden on the model to infer cross-sectional structure from raw returns alone.
- **Multi-Horizon Momentum (mom_13w, mom_26w, mom_52w):** Cumulative returns over trailing 13-, 26-, and 52-week windows (three, six, and twelve months) are included for every variate. The 52-week term is the classic Jegadeesh-Titman momentum factor², among the most robust signals in cross-sectional asset pricing. Whereas the 4-week RSI and price-versus-SMA features capture only short-horizon directional strength, the multi-horizon momentum features give the Variable Selection Network explicit medium-term trend information — the 3-to-12-month regimes that drive sector rotation. Adding these three features lifted the out-of-sample walk-forward Sharpe materially (Section V) while reducing the validation-to-test generalization gap, indicating that the added signal generalizes rather than being memorized.

All features are computed independently for SPY and for each of the eleven sector ETFs, then z-score normalized using statistics computed strictly from the training portion of each fold to prevent look-ahead leakage. Missing values are tolerated up to a maximum ratio of 5% per ticker; tickers exceeding this threshold are excluded from the relevant fold.

² Jegadeesh, Narasimhan, and Sheridan Titman. "Returns to Buying Winners and Selling Losers: Implications for Stock Market Efficiency." *The Journal of Finance*, vol. 48, no. 1, 1993, pp. 65–91.

III. Objective Function: The Ranking Advantage

The choice of objective function is a key methodological decision in this study, and the project evolved through two formulations. The model was initially trained with Mean Squared Error (MSE) on predicted vs. actual relative returns. MSE is the standard choice for regression problems and is the most intuitive loss for a model that emits numerical predictions. However, examination of the training dynamics revealed a pathology specific to the small-magnitude nature of weekly relative returns: MSE collapsed trivially toward zero (training loss fell by a factor of 600 within 100 epochs) while validation rank correlation crawled to only 0.115. The optimizer was driven primarily by predicting numbers near zero, providing weak gradient signal for the cross-sectional ranking task that the strategy actually depends on.

The model was subsequently retrained with a RankNet³-style pairwise ranking loss. For each pair of sectors (i, j) in each training week, the loss penalizes the model when the predicted ordering $\text{pred}[i]$ vs. $\text{pred}[j]$ disagrees with the actual ordering $\text{target}[i]$ vs. $\text{target}[j]$:

$$L_{\text{pair}} = \text{softplus}(-(\text{pred}[i] - \text{pred}[j]) \cdot \text{sign}(\text{target}[i] - \text{target}[j]))$$

Pairs with effectively-tied targets are excluded from the loss. The implementation is fully vectorized, producing all valid pairs per batch in a single tensor operation. This loss directly optimizes for the ranking quality that the long/short strategy depends on, rather than for absolute return accuracy.

In a sector rotation strategy, the absolute predicted return is less critical than the accuracy of the ranking. By maximizing the correlation between predicted and actual sector ranks, the model is incentivized to correctly identify the leaders and laggards, which is the mechanical requirement for a long/short strategy. The empirical impact of this change is documented in Section V.

IV. Strategy Execution & Walk-Forward Simulation

To ensure the findings are not artifacts of overfitting, we utilize a Walk-Forward validation cycle. This simulates a real-world trading environment by ensuring the model never sees data from the future during training.

- **Training Routine:** The model is trained on an 8-year rolling window (~400 weeks) and validated on a 1-year period, followed by testing on the subsequent 6 months (~26 weeks).

³ Burges, C., Shaked, T., Renshaw, E., Lazier, A., Deeds, M., Hamilton, N., & Hullender, G. (2005). Learning to rank using gradient descent. *Proceedings of the 22nd International Conference on Machine Learning - ICML '05*, 89-96. <https://doi.org/10.1145/1102351.1102363>

- **Market-Neutral Posture:** The strategy goes long the top 3 predicted outperformers and short the bottom 3 underperformers.
- **Ensembling:** Predictions are averaged across multiple random seeds (e.g., 5 seeds per fold) to stabilize the attention weights.




Figure: 10-seed ensemble walk-forward of the long/short sector rotation strategy (hybrid loss, with multi-horizon momentum features) across 36 folds spanning 504 weeks (approximately ten years out-of-sample). Top: cumulative return reaching approximately 2.1 at an aggregate Sharpe ratio of 1.82 and a 63% weekly win rate. Middle: per-fold Sharpe ratios (green positive, red negative), 28 of 36 folds positive, median fold Sharpe 1.76. Bottom: per-sector prediction correlation aggregated across all folds, strongest for XLRE and XLC.

V. The Sharpe Ratio: Quantifying Risk-Adjusted Performance

For a market-neutral sector rotation strategy, the Sharpe ratio serves as the primary performance metric. It measures the magnitude of excess returns relative to their volatility, which is the relevant question for a strategy designed to generate steady, uncorrelated alpha rather than to amplify market beta.

- **Mathematical Form:** $\text{Sharpe} = (\text{mean return} - \text{risk-free rate}) / \text{standard deviation of returns}$. For weekly returns, the value is annualized by multiplication by $\sqrt{52}$, giving comparability with strategies operating at other frequencies.

- **Why It Suits This Project:** Because the long/short construction removes broad-market exposure, the strategy is expected to produce a modest but persistent return stream. The Sharpe ratio rewards exactly this profile—consistency over magnitude—and is the metric institutional allocators expect to see for any market-neutral strategy.
- **Interpretation Bands:** Sharpe below 0.5 is typically considered weak or indistinguishable from noise; 0.5 to 1.0 reflects the historical range of major equity risk factors after costs; 1.0 to 2.0 covers most successful systematic equity strategies; values sustained above 3.0 on meaningful capital are vanishingly rare and warrant skepticism.
- **Single-Seed vs. Ensemble Reporting:** A single-seed walk-forward Sharpe is biased upward as an estimator of expected out-of-sample performance, since selecting the best result among multiple random initializations does not generalize to future runs. The ensemble mean across multiple seeds provides the more honest expected value, while the standard deviation across seeds quantifies the uncertainty around that estimate.
- **Limitations:** The Sharpe ratio assumes returns are approximately normally distributed and penalizes upside volatility identically to downside volatility. The Sortino ratio addresses the latter concern by penalizing only downside variance. Despite these limitations, the Sharpe ratio remains the de facto standard in quantitative finance for risk-adjusted performance evaluation, and is reported alongside cumulative return and maximum drawdown in any rigorous strategy evaluation.

Following correction of the loss function (transition from MSE to pairwise ranking loss, described in Section III) and a corresponding fix to model checkpoint selection (saving the model that maximized validation rank correlation rather than validation loss), a fresh single-training-run evaluation was performed using a 70/15/15 train/validation/test split. This evaluation served as a verifiable baseline benchmark for the corrected methodology. Three configurations were evaluated as a controlled ablation:

- **MSE baseline:** test rank correlation 0.095, test top-3 accuracy 34.8%. Best epoch 91 of 100 (training had not yet converged at termination).
- **Pairwise ranking loss:** test rank correlation 0.138, test top-3 accuracy 36.0%. A 45% relative improvement in rank correlation over the MSE baseline.
- **Pairwise loss with rank-correlation-based checkpoint selection:** test rank correlation 0.170, test top-3 accuracy 34.5%. An additional 23% relative improvement over the previous configuration, and 79% total improvement over the MSE baseline.

The progression demonstrates that each methodological refinement contributed measurable improvement, with the loss-function change being the largest single

contributor. The final test rank correlation of 0.170 corresponds, by proportional scaling from the MSE-baseline strategy Sharpe of approximately 0.50 observed on the same evaluation split, to an expected Sharpe ratio in the range of 0.75 to 1.0—a controlled single-split estimate that motivated the full ensemble walk-forward reported below. Per-sector predictive correlation in this evaluation showed positive signal on 7 of 11 sectors, with the strongest contributors being XLC (correlation 0.31) and XLRE (0.29). The broader distribution of predictive signal across sectors, relative to the XLRE-concentrated pattern observed in the preliminary walk-forward, suggests the corrected training pipeline is extracting genuine cross-sectional ranking signal rather than relying disproportionately on a single sector.

Definitive Evaluation

The definitive evaluation combines all methodological improvements—hybrid pairwise-plus-MSE loss for both ranking quality and scale calibration, macroeconomic features (Treasury, credit, dollar, and volatility variates added alongside the eleven sector ETFs), multi-horizon momentum features (described in Section II), rank-correlation-based checkpoint selection, and increased dropout regularization—in a 10-seed ensemble walk-forward spanning 36 folds and 504 weeks, roughly ten years of strictly out-of-sample testing. This evaluation produced an aggregate Sharpe ratio of 1.82 with a 63% weekly win rate; approximately 78 percent of folds (28 of 36) were positive. Per-fold Sharpe ratios ranged from roughly -1.9 to $+8.0$, with a median fold Sharpe of 1.76. The cumulative return curve climbed steadily from zero to approximately 2.1 over the full period, traversing multiple distinct market regimes—including the 2018 selloff, the 2020 pandemic crash, the 2022 bear market, and the 2023–2024 recovery—without extended drawdowns.

This result reflects two distinct improvements over the prior decade-spanning evaluation, which produced a Sharpe of 1.16 without multi-horizon momentum features. First, adding the momentum features (Section II) raised the single-seed walk-forward Sharpe from 1.16 to 1.60, with the mean per-fold prediction correlation improving in parallel; a controlled ablation confirmed the gain was attributable to the features rather than to any change in the loss function. Second, expanding from a single seed to a 10-seed ensemble raised the Sharpe further to 1.82 and the mean per-fold prediction correlation from 0.154 to 0.197. That the ensemble improved upon the single-seed result, rather than regressing toward it, indicates that single-seed runs were understating the model’s signal through initialization noise rather than overstating it through favorable initialization—evidence that the edge is structural rather than an artifact of a lucky seed. The ten-year, 36-fold sample provides far narrower confidence bounds than any small-sample evaluation, and the steady equity curve across heterogeneous regimes indicates a persistent structural edge rather than a small number of favorable folds. The per-sector predictive correlations, while lower in magnitude than single-period estimates (XLRE 0.16, XLC 0.11), reflect

the honest long-run signal strength after averaging across a decade of varied conditions; a single fold may exhibit a much stronger XLRE correlation, but the decade-averaged figure is the appropriate basis for forward expectations. Accounting for the anticipated 15 to 25 percent degradation from real-world frictions—transaction costs, slippage, and option bid-ask spreads—a realistic deployed Sharpe ratio is estimated in the range of 0.9 to 1.0, which remains a solid risk-adjusted return for a market-neutral strategy.

This 1.82 decade-spanning ensemble result is the headline performance characterization of the strategy. It is the most reproducible and least sample-dependent number available. It should be read as the Sharpe of the underlying long/short sector signal under idealized execution; the option-based implementation described in Section VII introduces additional friction and convexity, and after the anticipated 15 to 25 percent real-world degradation a realistic deployed Sharpe is estimated in the range of approximately 1.4 to 1.55—still a strong risk-adjusted return for a market-neutral strategy, and substantially improved over the pre-momentum baseline.

VI. Model Interpretability: Variable Selection

A central design feature of the Temporal Fusion Transformer is interpretability through its Variable Selection Network. At every timestep the network emits a softmax distribution over the input variates, and because these weights are normalized to sum to one, they can be read directly as the proportion of the model’s attention allocated to each input. Averaging these weights across all samples and timesteps yields a global measure of which inputs the model relies upon. The analysis below was produced by capturing the variable selection weights from the trained model across the full evaluation history.




Figure: Average Variable Selection Network weight per input variate. Green bars are sector ETFs (prediction targets), blue bars are macroeconomic context variates, gray is the SPY benchmark. The dashed line marks equal weight (5.6%).

The complete ranking underlying the chart is given in the table below, ordered by weight and annotated with each variate's group and a proportional bar. The group totals at the foot of the table summarize the headline finding: macroeconomic context variates collectively command the majority of the model's attention.

=====				
VARIABLE IMPORTANCE (avg VSN weight, all samples)				
=====				
Rank	Ticker	Group	Weight	

1	IEF	macro	15.58%	#####
2	UUP	macro	14.38%	#####
3	XLI	sector	11.44%	#####
4	VXX	macro	10.61%	#####
5	XLY	sector	8.02%	#####
6	XLE	sector	7.63%	#####
7	HYG	macro	7.33%	#####
8	XLP	sector	4.93%	#####
9	SPY	benchmark	4.60%	#####
10	TLT	macro	3.44%	#####
11	XLB	sector	2.29%	#####
12	LQD	macro	2.24%	#####
13	XLV	sector	1.87%	#####
14	XLF	sector	1.44%	#####
15	XLRE	sector	1.40%	#####
16	XLU	sector	1.39%	#####
17	XLK	sector	0.84%	#####
18	XLC	sector	0.58%	#####

Group totals: sectors 41.8% macros 53.6% benchmark 4.6%				
=====				

Table: Complete variable selection weights for all eighteen input variates, ranked by average attention. Intermediate Treasuries (IEF), the US dollar (UUP), and equity volatility (VXX) occupy three of the top four positions; the sectors with the strongest walk-forward predictive correlation (XLRE, XLC) rank near the bottom in direct weight, indicating they are forecast from macroeconomic context rather than their own price histories.

The result is striking: the model allocates the majority of its attention to macroeconomic context rather than to the sector price histories it ultimately ranks. The six macro variates collectively receive 53.6 percent of total weight, the eleven sectors receive 41.8 percent, and the SPY benchmark receives 4.6 percent. The three highest-weighted individual inputs in the entire model are intermediate Treasuries (IEF, 15.58 percent), the US dollar (UUP, 14.38 percent), and equity volatility (VXX, 10.61 percent); together with high-yield credit (HYG, 7.33 percent), these four macro signals account for roughly 48 percent of the model's attention.

This reveals an economically coherent strategy that the model discovered on its own. It reads the macroeconomic regime—the level of intermediate rates, the strength of the dollar, the state of equity volatility, and the condition of credit markets—and uses a small number of bellwether cyclical sectors as confirmation, principally Industrials (XLI, 11.44 percent), Consumer Discretionary (XLY, 8.02 percent), and Energy (XLE, 7.63 percent). It then ranks all eleven sectors off that macro read. This is closely analogous to how a

discretionary macro strategist would approach sector rotation, and the Variable Selection Network arrived at the structure without explicit instruction.

A notable inversion supports this interpretation. The sectors that exhibited the strongest predictive correlation in walk-forward evaluation—Real Estate (XLRE) and Communication Services (XLC)—received the lowest direct variable selection weight, at 1.40 percent and 0.58 percent respectively. The model is therefore not forecasting these sectors from their own price histories; it is forecasting them from the shared macroeconomic drivers. Rate-sensitive and defensive sectors in particular appear to be predicted indirectly through the macro variates rather than read directly, which is consistent with the economic reality that such sectors are largely driven by the rate and risk environment.

The decomposition also exposes redundancy that informs future model design. Within the rates complex, intermediate Treasuries (IEF) dominate long Treasuries (TLT, 3.44 percent) by more than fourfold, indicating that the belly of the yield curve carries substantially more sector-rotation signal than the long end. Within credit, high-yield (HYG) dominates investment-grade (LQD, 2.24 percent). A leaner macro set retaining only IEF, UUP, VXX, and HYG would likely preserve most of the model's informational advantage while reducing parameter count and concentration. Conversely, the heavy reliance on six macro series—over half of total attention—represents a concentration that warrants monitoring, as the strategy is implicitly exposed to those series behaving in the future as they did historically.

VII. Real-World Friction & Option Modeling

The strategy execution involves a simulation engine that accounts for variables often ignored in academic backtests, bridging the gap between theoretical alpha and tradable returns.

- **Derivative Implementation:** The model simulates buying \$1.00 OTM calls for long picks and OTM puts for shorts, with roughly 30 days to expiration.
- **Friction and Decay:** The simulation includes per-contract commissions and accounts for time decay (Theta) and Bid/Ask spreads.

VIII. Conclusion

The Temporal Fusion Transformer provides a resilient signal for sector rotation by focusing on cross-sectional ranking and utilizing an interpretable architecture. Future research will focus on incorporating alternative data sets.

Appendix: Command-Line Interface and Execution Flags

1. Core Pipeline (run.py)

Main entry point to execute the full end-to-end lifecycle.

- **--skip-download:** Bypasses yfinance data retrieval.
- **--skip-preprocess:** Skips feature engineering and loads existing tensors.
- **--epochs [INT]:** Overrides default training duration.
- **--device [STR]:** Forces execution on a specific device (cpu, cuda, mps).
- **--tune:** Enables Optuna hyperparameter search prior to final training.

2. Backtesting and Validation (walk_forward.py)

- **--train-weeks [INT]:** Weeks in the rolling training window (default 400).
- **--val-weeks [INT]:** Size of the validation window (default 52).
- **--test-weeks [INT]:** Size of the out-of-sample test window (default 26).
- **--n-seeds [INT]:** Model initializations to ensemble per fold.

3. Strategy Simulation (simulation.py)

Command: setup

- **--capital [FLOAT]:** Starting cash balance.
- **--otm [FLOAT]:** Strike distance from current price.
- **--commissions [FLOAT]:** Fee per contract per side.

Command: predict, execute, status, close, week, reset